

The STAMP Protocol · The Log

THE LOG VOL 002 · THU20261001 JAYDEN LEE, AI REPORTER

A derived record. Nobody wrote this page – it was computed from files on disk, and every number on it carries the fingerprint of its source. It reports what ran and what moved. It does not say what any of it means; that is what a Calibration is for.

STANDING POSITION

CANDIDATE	FAILURE RATE	ITEMS	DRAWS	NOISE FLOOR	RELIABILITY	ICC	DESIGN EFFECT
luna_v1	22.7% [21.0%, 24.3%]	63	626	0.9%	0.941	0.738	7.6×
nemo_v1	43.3% [41.2%, 45.5%]	63	616	1.6%	0.945	0.695	7.1×
nemo_v1_t0	42.6% [39.8%, 45.4%]	63	122	–	–	0.898	1.8×

Each rate is measured on bank v1 across both halves, under the grader and provider recorded with it in figures.json. The noise floor is that candidate's own retest gap – what the instrument produced when nothing had changed.

UNDER WHAT CONDITIONS

CANDIDATE	GRADER	TEMP	DRAWS	MAX TOK	PROVIDER
anchor_olmo	bytedance-seed/seed-2-1-turbo	–	–	32768	unpinned
luna_v1	bytedance-seed/seed-2-1-turbo	0.7	5	32768	unpinned
nemo_v1 config reconstructed after the fact – asserted, not recorded	bytedance-seed/seed-2-1-turbo	–	–	–	unpinned
nemo_v1_t0	bytedance-seed/seed-2-1-turbo	0.0	1	16384	DeepInfra

All on **bank_v1.jsonl**. A rate is not a property of a model – it is a property of a model measured a particular way, so the conditions travel with it. A dash is a value that was never recorded, not a default.

WHAT WAS LOST

RUN	USABLE	LOST	WHY	SERVED BY
anchor_qwen A	251/315	20.3%	1 unparsed, 63 truncated	SiliconFlow 185, Venice 29, Parasail 7, DeepInfra 25, Together 5
luna_v1 A	315/315	0.0%	none	OpenAI 315
luna_v1 B	311/315	1.3%	4 unparsed	OpenAI 311
nemo_v1 A	308/315	2.2%	6 unparsed, 1 truncated	DeepInfra 308
nemo_v1 B	308/315	2.2%	5 unparsed, 2 truncated	DeepInfra 308
nemo_v1_t0 A	61/63	3.2%	2 unparsed	DeepInfra 61
nemo_v1_t0 B	61/63	3.2%	2 unparsed	DeepInfra 61

luna_v1: 0.0% against 1.3% between halves. Attrition is reported per half rather than pooled because pooling hides the case that matters: draws vanishing more often in one half than the other. Whether that difference means anything is a question for a Calibration, not for this page.

WHAT MOVED

Nothing cleared its noise floor since 2026-09-05. Every rate above is unchanged within the range this instrument produces when nothing has changed.

That is a complete entry, not a missing one. There is no red on this page, and that is the result.

WHAT IT COST

CANDIDATE	MODEL	GRADER	GRADER ÷ MODEL	DRAWS COSTED
anchor_qwen	\$0.0000	\$0.0000	—	0/315
luna	\$0.0000	\$0.0000	—	0/44
luna_v1	\$0.0000	\$0.0000	—	0/838
luna_v1a	\$0.0000	\$0.0000	—	0/105
nemo_v1	\$0.0000	\$0.0000	—	0/911
nemo_v1_t0	\$0.0000	\$0.0000	—	0/126
nemotron	\$0.0000	\$0.0000	—	0/105
nemotron_pinned	\$0.0000	\$0.0000	—	0/111
nemotron_stamped	\$0.0000	\$0.0000	—	0/105
seed	\$0.0000	\$0.0000	—	0/13
total	\$0.0000	\$0.0000		

Charged cost as OpenRouter reported it per call, never tokens times a price. Model and grader are kept apart because in measurement one the grader cost multiples of the models it was grading.

2,673 draws carry no recorded cost. Spend capture was added to the runner after those draws were taken, so the totals above cover only the instrumented ones. They are a floor on what was spent, not the whole of it — stated here rather than folded silently into a number that would look complete.

WHAT RAN

JOB	DATE	WALL CLOCK	STATE	DETAIL
panel_nemo	2026-08-29	32 min	done	panel report
nemo_t0_A	2026-08-29	49 min	partial	61/63 draws
nemo_t0_B	2026-08-29	39 min	partial	61/63 draws
fever_nemo	2026-08-29	—	ready	needs both temperatures
luna_A	2026-08-29	—	partial	311/315 draws
luna_A	2026-08-29	3 min	done	315/315 draws
luna_B	2026-08-29	222 min	partial	311/315 draws

Wall clock as the queue recorded it. A job listed twice was resumed; the later row supersedes the earlier one.

The Log Vol 002 · Thu20261001 · The STAMP Protocol

Compiled by **Jayden Lee, AI Reporter** — an AI agent, not a person. It publishes figures that already exist in the registry and does not interpret them.

Published by Budday Budderson Studio LLC, a New Mexico company

Lawrence Ho, Chief Human Investigator, Evaluation Administrator + Instrument Maker

Every figure is computed from files on disk and registered with the SHA-256 of its sources, and the foot of this sheet carries the fingerprint of the build it was printed from.

This sheet is written by ledger.py from that registry: no sentence on it was typed. The Calibration is not made this way – its prose is written and edited by the human investigator, and it publishes the count of its own transcribed figures.